
EMBER[®] APPLICATION DEVELOPMENT FUNDAMENTALS: TOKENS

This document describes tokens and shows how to use them in code. It also discusses bindings, the application-defined associations between two devices on a network.

New in This Revision

Initial release (contents previously published).

Contents

1	Introduction	2
1.1	Purpose	2
1.2	Usage	3
2	Accessing Standard (Non-indexed) Tokens	3
3	Accessing Indexed Tokens	4
4	Accessing Counter Tokens	4
5	Accessing Manufacturing Tokens	4
6	Custom Tokens	5
6.1	Mechanics	5
6.1.1	Define the Token Name	5
6.1.2	Define the Token Type	6
6.1.3	Define the Token Storage	6
6.2	Default Tokens	7
7	Bindings	7
8	For More Information	7

UG103.7

A token is an abstract data constant that has special persistent meaning for an application. This document describes tokens and shows how to use them in code. A token can be one of several types:

- Standard
- Indexed
- Counter
- Manufacturing

Tokens can be categorized in the following general groups:

- Stack tokens – these tokens are read/write and defined in every application to support stack behavior. These tokens live in Simulated EEPROM.
- Application tokens – these tokens are read/write and defined by the application to support application behavior. It is left open to a customer to decide if there are application tokens, how they are defined, and what they do. These tokens live in Simulated EEPROM.
- Fixed manufacturing tokens – these tokens are non-modifiable read only tokens specific to each individual chip and set during chip production. They support key low-level and high-level functionality.
- Customer manufacturing tokens – these tokens are read only from on chip code and read/write from external programming tools. Some of these tokens have a defined purposes to support stack behavior. It is left open to a customer to decide if there are any additional manufacturing tokens, how they are defined, and what they do.

The document also discusses bindings, the application-defined associations between two devices on a network.

1 Introduction

A token has two parts: a token key and token data. The token key is a unique identifier that is used to store and retrieve the token data. In many cases, the word "token" is used quite loosely to mean the token key, the token data, or the combination of key and data. Usually it is clear from the context which meaning to use. In this document token always refers to the key + data pair.

Tokens are typically stored in NVRAM which is most commonly directly in flash (manufacturing tokens) or in the Simulated EEPROM (stack and application tokens) which operates on top of flash.

1.1 Purpose

The fundamental purpose of the token system as compared to generic RAM usage is to allow the token data to persist across reboots and during power loss. By using the token key to identify the proper data, the application requesting the token data does not need to know the exact storage location of the data. This simplifies application design and code reuse. Tokens are also useful when the underlying storage of the data may change over time across implementations.

The EM2xx and EM3xx chips use a special memory-rotation algorithm, called Simulated EEPROM, to prevent premature overuse of the underlying flash. Ember software reserves a section of flash for use by the Simulated EEPROM. EM2xx chips reserve 8 kB of flash. EM3xx chips have the choice of two Simulated EEPROM sizes: 4 kB, 8 kB (the default). The Simulated EEPROM storage area implements a wear-leveling algorithm to extend the number of Simulated EEPROM write cycles beyond the physical limit of the flash. Silicon Labs recommends that application designers familiarize themselves with the Simulated EEPROM and its limitations, so that they design the application's use of tokens for optimal flash write cycles. Refer to document AN703, *Using the Simulated EEPROM*, for more information. However, application designers do not need to know any of the underlying details. Understanding how to use token keys is enough.

1.2 Usage

EmberZNet PRO provides a simple set of APIs for accessing token data. The full documentation may be found in EmberZNet PRO API reference for your platform (120-3022-000 for the EM35x, 120-3016-000 for the EM250, and 120-3020-000 for the EM260).

The basic API functions include:

```
void halCommonGetToken( data, token )
void halCommonSetToken( token, data )
```

In this case, 'token' is the token key, and 'data' is the token data.

Two special types of tokens are available: **indexed tokens** and **counter tokens**. An indexed token can be used when the data to be stored is an array in which each element may be accessed and updated independently from the others. A counter token can be used when the default operation on the token is to retrieve it, increment it by one, and then store it again.

The API functions for these two types of tokens are:

```
void halCommonGetIndexedToken( data, token, index )
void halCommonSetIndexedToken( token, index, data )
void halCommonIncrementCounterToken( token )
```

The counter token can also be accessed with the normal `halCommonGetToken()`, `halCommonSetToken()` calls. Setting a counter token severely degrades Simulated EEPROM performance, and therefore should only be done when absolutely necessary.

2 Accessing Standard (Non-indexed) Tokens

Some applications may need to store configuration data at installation time. Usually, this is a good application for a standard token. Assume the token is defined to use the token key `DEVICE_INSTALL_DATA`, and the data structure looks like this:

```
typedef struct {
    int8u install_date[11] /** YYYY-mm-dd + NULL */
    int8u room_number; /** The room where this device is installed */
} InstallationData_t;
```

Then you can access it with a code snippet like this:

```
InstallationData_t data;
// Read the stored token data
halCommonGetToken(&data, TOKEN_DEVICE_INSTALL_DATA);
// Set the local copy of the data to new values
data.room_number = < user input data >
MEMCOPY(data.install_date, < user input data>, 0, sizeof(data.install_date));
// Update the stored token data with the new values
halCommonSetToken(TOKEN_DEVICE_INSTALL_DATA, &data);
```

Instructions for configuring custom tokens are provided in Section 6, Custom Tokens.

3 Accessing Indexed Tokens

To store a set of similar values, such as an array of preferred temperature settings throughout the day, use the default data type `int16s` to store the desired temperatures, and define a token called `HOURLY_TEMPERATURES`.

A local copy of the entire data set would look like this:

```
int16s hourlyTemperatures[HOURLY_IN_DAY]; /** 24 hours per day */
```

In the application code, you can access or update just one of the values in the day using the indexed token functions:

```
int16s getCurrentTargetTemperature(int8u hour) {
    int16s temperatureThisHour = 0; /** Stores the temperature for return */
    if (hour < HOURLY_IN_DAY) {
        halCommonGetIndexedToken(&temperatureThisHour,
            TOKEN_HOURLY_TEMPERATURES, hour);
    }
    return temperatureThisHour;
}

void setTargetTemperature(int8u hour, int16s targetTemperature) {
    if (hour < HOURLY_IN_DAY) {
        halCommonSetIndexedToken(TOKEN_HOURLY_TEMPERATURE, hour,
            &temperatureThisHour);
    }
}
```

4 Accessing Counter Tokens

Counting the number of heating cycles a thermostat has initiated is a perfect use for a counter token. Assume it is named `LIFETIME_HEAT_CYCLES`, and it is an `int32u`.

```
void requestHeatCycle(void) {
    /// < application logic to initiate heat cycle >
    halCommonIncrementCounterToken(TOKEN_LIFETIME_HEAT_CYCLES);
}

int32u totalHeatCycles(void) {
    int32u heatCycles;
    halCommonGetToken(&heatCycles, TOKEN_LIFETIME_HEAT_CYCLES);
    return heatCycles;
}
```

5 Accessing Manufacturing Tokens

Manufacturing tokens are defined and treated as standard (non-indexed) tokens. The major difference is manufacturing tokens can only be read from on chip code and customer manufacturing tokens can only be written with external programming tools. In addition to being able to read manufacturing tokens with the usual standard token read, `halCommonGetToken()`, manufacturing tokens can also be read with their own dedicated API, `halCommonGetMfgToken()`. This `MfgToken` API takes the same parameters as the standard `GetToken` API. In general, MFG tokens should be accessed through the standard `GetToken` API. The two primary purposes for using the `MfgToken` API is for slightly faster access and early on in the boot process before `emberInit()` is called.

6 Custom Tokens

Custom application tokens are defined in a header file. The header file can be specific to each project, and is defined by the preprocessor variable `APPLICATION_TOKEN_HEADER`. Custom manufacturing tokens require editing an existing header file. For the EM35x, this header file is located at `hal/micro/cortexm3/token-manufacturing.h`. For the EM250, this header is located at `hal/micro/avr-atmega/token-manufacturing.h`. Defining custom manufacturing tokens is generally considered a rare practice, but is still possible. Note that to define them requires editing a file that contains key definitions critical to the system so making any additions should be done very carefully and should follow the same format as the other definitions in the file.

The `APPLICATION_TOKEN_HEADER` file should have the following structure:

```
/**
 * Custom Application Tokens
 */
// Define token names here
#ifdef DEFINETYPES
// Include or define any typedef for tokens here
#endif //DEFINETYPES
#ifdef DEFINETOKENS
// Define the actual token storage information here
#endif //DEFINETOKENS
```

Note: The header files do not have `#ifndef HEADER_FILE / #define HEADER_FILE` sequence at the top. This is important because this header file is included several times for different purposes.

6.1 Mechanics

Adding a custom token to the header file involves three steps:

1. Define the token name.
2. Add any typedef needed for the token, if it is using an application-defined type.
3. Define the token storage.

The following illustrates how to define the three previous token examples.

6.1.1 Define the Token Name

When defining the name, do not prepend the word `TOKEN`. Instead, use the word `CREATOR`:

```
/**
 * Custom Application Tokens
 */
// Define token names here
#define CREATOR_DEVICE_INSTALL_DATA (0x000A)
#define CREATOR_HOURLY_TEMPERATURES (0x000B)
#define CREATOR_LIFETIME_HEAT_CYCLES (0x000C)
```

This defines the token key and links it to a programmatic variable. The token names are actually `DEVICE_INSTALL_DATA`, `HOURLY_TEMPERATURES`, and `LIFETIME_HEAT_CYCLES`, with different tags prepended to the beginning depending on the usage. Thus they are referred to in the example code as `TOKEN_DEVICE_INSTALL_DATA`, and so on.

The token key is a 16-bit value that must be unique within this device. The first-bit is reserved for manufacturing and stack tokens, so all custom tokens should have a token key less than 0x8000.

UG103.7

The token key is critical to linking application usage with the proper data and as such a unique key should always be used when defining a new token or even changing the structure of an existing token. Always using a unique key guarantees a proper link between application and data.

6.1.2 Define the Token Type

Each token in this example is a different type; however, the `HOURLY_TEMPERATURES` and `LIFETIME_HEAT_CYCLES` types are built-in types in C. Only the `DEVICE_INSTALL_DATA` type is a custom data structure.

The token type is defined using the structure introduced at the start of section 6, Custom Tokens, and repeated next. Note that the token type must be defined only in one place, as the compiler will complain if the same data structure is defined twice.

```
#ifndef DEFINETYPES
// Include or define any typedef for tokens here
typedef struct {
    int8u install_date[11] /** YYYY-mm-dd + NULL */
    int8u room_number; /** The room where this device is installed */
} InstallationData_t;
#endif //DEFINETYPES
```

6.1.3 Define the Token Storage

After any custom types are defined, the token storage is defined. This informs the token management software about the tokens being defined. Each token, whether custom or built-in, gets its own entry in this part:

```
#ifndef DEFINETOKENS
// Define the actual token storage information here
DEFINE_BASIC_TOKEN(DEVICE_INSTALL_DATA,
    InstallationData_t,
    {0, {0,...}})
DEFINE_INDEXED_TOKEN(HOURLY_TEMPERATURES, int16u, HOURS_IN_DAY, {0,...})
DEFINE_COUNTER_TOKEN(LIFETIME_HEAT_CYCLES, int32u, 0)
#endif //DEFINETOKENS
```

The following expands on each step in this process.

```
DEFINE_BASIC_TOKEN(DEVICE_INSTALL_DATA,
    InstallationData_t,
    {0, {0,...}})
```

`DEFINE_BASIC_TOKEN` takes three arguments: the name (`DEVICE_INSTALL_DATA`), the data type (`InstallationData_t`), and the default value of the token if it has never been written by the application (`{0, {0,...}}`).

The default value takes the same syntax as C default initializers. In this case, the first value (`room_number`) is initialized to 0, and the next value (`installation_date`) is set to all 0s because the `{0,...}` syntax fills the remainder of the array with 0.

The syntax of `DEFINE_COUNTER_TOKEN` is identical to `DEFINE_BASIC_TOKEN`.

`DEFINE_INDEXED_TOKEN` requires a length of the array -- in this case, `HOURS_IN_DAY`, or 24. Its final argument is the default value of every element in the array. Again, in this case it is initialized to all 0.

The next section discusses default tokens, including stack tokens and manufacturing tokens.

6.2 Default Tokens

The EmberZNet PRO stack contains some default tokens that may be useful for the application developer. These tokens come in two flavors:

Stack Tokens are runtime configuration options set by the stack; these should not be changed by the application.

Manufacturing Tokens are set at manufacturing time and cannot be changed by the application.

To view the stack tokens, refer to the file:

```
<install-dir>/stack/config/token-stack.h
```

To view the manufacturing tokens for your chip, refer to one of the follow files:

```
<install-dir>/hal/micro/xap2b/token-manufacturing.h
```

```
<install-dir>/hal/micro/cortexm3/token-manufacturing.h
```

Search for `CREATOR` to see the defined names. If the entire file seems overwhelming, focus only on the section describing the tokens.

Some of the fixed manufacturing tokens may be set by the manufacturer when the board is created. For example, a custom EUI-64 address may be set by the vendor to override the internal EUI-64 address provided by Silicon Labs. Other tokens, such as the internal EUI-64, cannot be overwritten.

For more information about manufacturing and token programming, refer to documents AN702, *Bringing Up Custom Devices for the EM250 SoC Platform*, AN704, *Bringing Up Custom Devices for the EM260 Co-Processor*, and AN710, *Bringing Up Custom Devices for the EM35x SoC Platform*.

7 Bindings

Bindings are application-defined associations between two devices on the network. A binding may be used when sending to any message destination (multicast group or unicast point). They are used to provide a persistent means to store a pairing of two devices and the clusters, endpoints, and destinations they are using. The binding table consumes 2 bytes of RAM and 13 bytes of flash for each entry. Parts of this table reside in RAM, but most of the data is stored in token data. Therefore, the number of available bindings is limited by the processor's available RAM and EEPROM. It is important to understand that your application is responsible for managing binding table entries. For detailed information on `EMBER_BINDING_TABLE_SIZE` and `TOKEN_BINDING_TABLE`, refer to the EmberZNet PRO API reference for your platform (120-3022-000 for the EM35x, 120-3016-000 for the EM250, and 120-3020 for the EM260).

The `EMBER_BINDING_TABLE_SIZE` constant specifies the maximum number of bindings supported by the stack. This includes the bindings in EEPROM and in RAM. It is set in `stack/include/ember-configuration-defaults.h`, and can be reset through preprocessor definitions at compile time.

8 For More Information

For more detailed information about the underlying token mechanism and the simulated EEPROM that provides the EM2xx and EM3xx NVRAM implementation, please refer to the EmberZNet PRO API reference for your platform (120-3022-000 for the EM35x, 120-3016-000 for the EM250, and 120-3020-000 for the EM260), which has detailed information about the use of these items.

CONTACT INFORMATION

Silicon Laboratories Inc.

400 West Cesar Chavez
Austin, TX 78701
Tel: 1+(512) 416-8500
Fax: 1+(512) 416-9669
Toll Free: 1+(877) 444-3032

Please visit the Silicon Labs Technical Support web page for ZigBee products:
www.silabs.com/zigbee-support and register to submit a technical support request

Patent Notice

Silicon Labs invests in research and development to help our customers differentiate in the market with innovative low-power, small size, analog-intensive mixed-signal solutions. Silicon Labs' extensive patent portfolio is a testament to our unique approach and world-class engineering team.

The information in this document is believed to be accurate in all respects at the time of publication but is subject to change without notice. Silicon Laboratories assumes no responsibility for errors and omissions, and disclaims responsibility for any consequences resulting from the use of information included herein. Additionally, Silicon Laboratories assumes no responsibility for the functioning of undescribed features or parameters. Silicon Laboratories reserves the right to make changes without further notice. Silicon Laboratories makes no warranty, representation or guarantee regarding the suitability of its products for any particular purpose, nor does Silicon Laboratories assume any liability arising out of the application or use of any product or circuit, and specifically disclaims any and all liability, including without limitation consequential or incidental damages. Silicon Laboratories products are not designed, intended, or authorized for use in applications intended to support or sustain life, or for any other application in which the failure of the Silicon Laboratories product could create a situation where personal injury or death may occur. Should Buyer purchase or use Silicon Laboratories products for any such unintended or unauthorized application, Buyer shall indemnify and hold Silicon Laboratories harmless against all claims and damages.

Silicon Laboratories, Silicon Labs, and Ember are registered trademarks of Silicon Laboratories Inc.

Other products or brandnames mentioned herein are trademarks or registered trademarks of their respective holders.